

Load the Data

```
In [1]: import pandas as pd
import zipfile
import os

# Function to load CSV from a ZIP file with multiple files
def load_csv_from_zip(zip_path, csv_filename):
    with zipfile.ZipFile(zip_path, 'r') as z:
        # Extract and read the specific CSV file
        with z.open(csv_filename) as f:
            return pd.read_csv(f)

# Define the relative path to the datasets folder
datasets_path = os.path.join '..', 'Datasets')

# Load datasets from zipped CSV files specifying the correct CSV filename
df_gb = load_csv_from_zip(os.path.join(datasets_path, 'GBvideos.csv.zip'), 'GBvideos.csv')
df_us = load_csv_from_zip(os.path.join(datasets_path, 'USvideos.csv.zip'), 'USvideos.csv')

# Merge 5 files into 1
merged_df = pd.concat([df_gb, df_us], ignore_index=True)

# Check the first few rows of the merged DataFrame
print(merged_df.head())
```

```

      video_id trending_date \
0  Jw1Y-zhQURU      17.14.11
1  3s1rvMFUweQ      17.14.11
2  n1WpP7iowLc      17.14.11
3  PUTEiSjKwJU      17.14.11
4  rHwDegptbI4      17.14.11

      title \
0  John Lewis Christmas Ad 2017 - #MozTheMonster
1  Taylor Swift: ...Ready for It? (Live) - SNL
2  Eminem - Walk On Water (Audio) ft. Beyoncé
3  Goals from Salford City vs Class of 92 and Fri...
4  Dashcam captures truck's near miss with child ...

      channel_title  category_id  publish_time
e \
0  John Lewis      26  2017-11-10T07:38:29.000
Z
1  Saturday Night Live      24  2017-11-12T06:24:44.000
7
```

Check Missing Values

```
In [2]: # Check for missing values in the merged DataFrame
print("Missing values")
print(merged_df.isnull().sum())
```

```
Missing values
video_id          0
trending_date     0
title             0
channel_title     0
category_id       0
publish_time      0
tags              0
views             0
likes             0
dislikes          0
comment_count     0
thumbnail_link    0
comments_disabled 0
ratings_disabled  0
video_error_or_removed 0
description       1182
dtype: int64
```

```
In [3]: df = merged_df.dropna()
```

```
In [4]: # Check for missing values in the merged DataFrame
print("Missing values")
print(df.isnull().sum())
```

```
Missing values
video_id          0
trending_date     0
title             0
channel_title     0
category_id       0
publish_time      0
tags              0
views             0
likes             0
dislikes          0
comment_count     0
thumbnail_link    0
comments_disabled 0
ratings_disabled  0
video_error_or_removed 0
description       0
dtype: int64
```

Exploratory Data Analysis (EDA)

Check Outliers

```

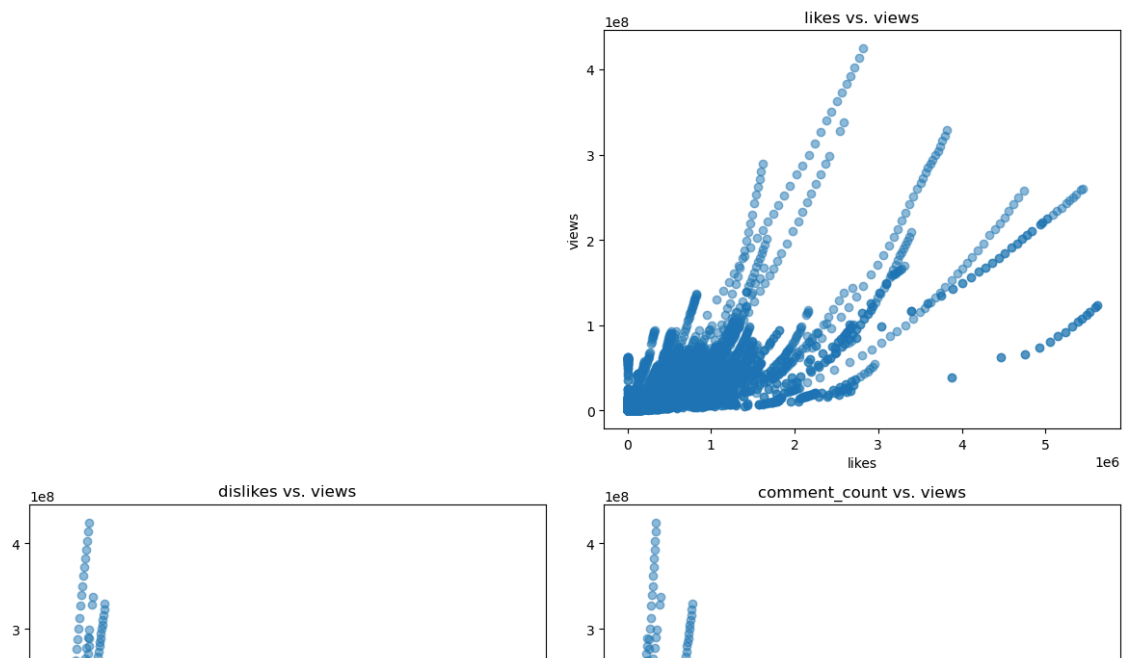
In [5]: import seaborn as sns
import matplotlib.pyplot as plt

# Define numerical columns
numerical_columns = ['views', 'likes', 'dislikes', 'comment_count']

# Scatter plots for each numerical column vs. 'views'
plt.figure(figsize=(12, 10))
for i, column in enumerate(numerical_columns, 1):
    if column != 'views':
        plt.subplot(2, 2, i)
        plt.scatter(merged_df[column], merged_df['views'], alpha=0.5)
        plt.title(f'{column} vs. views')
        plt.xlabel(column)
        plt.ylabel('views')

plt.tight_layout()
plt.show()

```



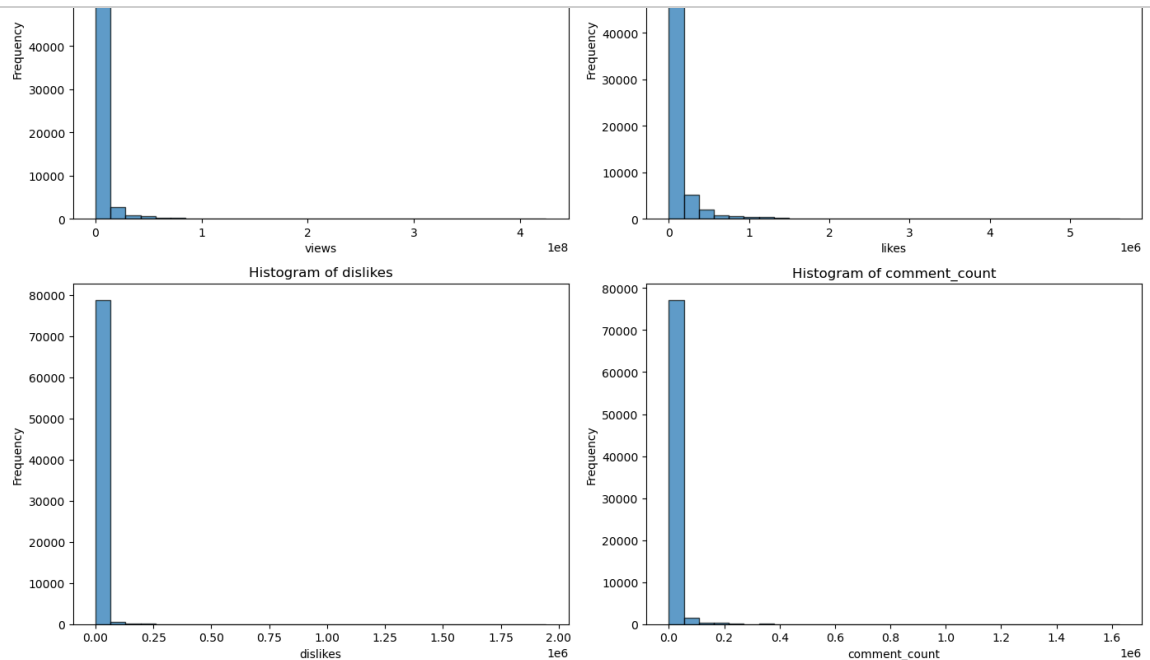
EDA for Numerical Variables

```
In [6]: #data exploration for numerical columns
import matplotlib.pyplot as plt

# Define numerical columns
numerical_columns = ['views', 'likes', 'dislikes', 'comment_count']

# Create histograms for each numerical column
plt.figure(figsize=(14, 10))
for i, column in enumerate(numerical_columns, 1):
    plt.subplot(2, 2, i)
    plt.hist(merged_df[column], bins=30, alpha=0.7, edgecolor='black')
    plt.title(f'Histogram of {column}')
    plt.xlabel(column)
    plt.ylabel('Frequency')

plt.tight_layout()
plt.show()
```




EDA for Boolean Variables

```
In [7]: import seaborn as sns

# Define boolean columns
boolean_columns = ['comments_disabled', 'ratings_disabled', 'video_error_or_removed']

# Plot bar plots for each boolean column
plt.figure(figsize=(14, 6))
for i, column in enumerate(boolean_columns, 1):
    plt.subplot(1, 3, i)
    # Count the occurrences of each boolean value
    counts = merged_df[column].value_counts()
    # Plot bar plot
    sns.barplot(x=counts.index, y=counts.values, palette='viridis')
    plt.title(f'Distribution of {column}')
    plt.xlabel('Value')
    plt.ylabel('Count')

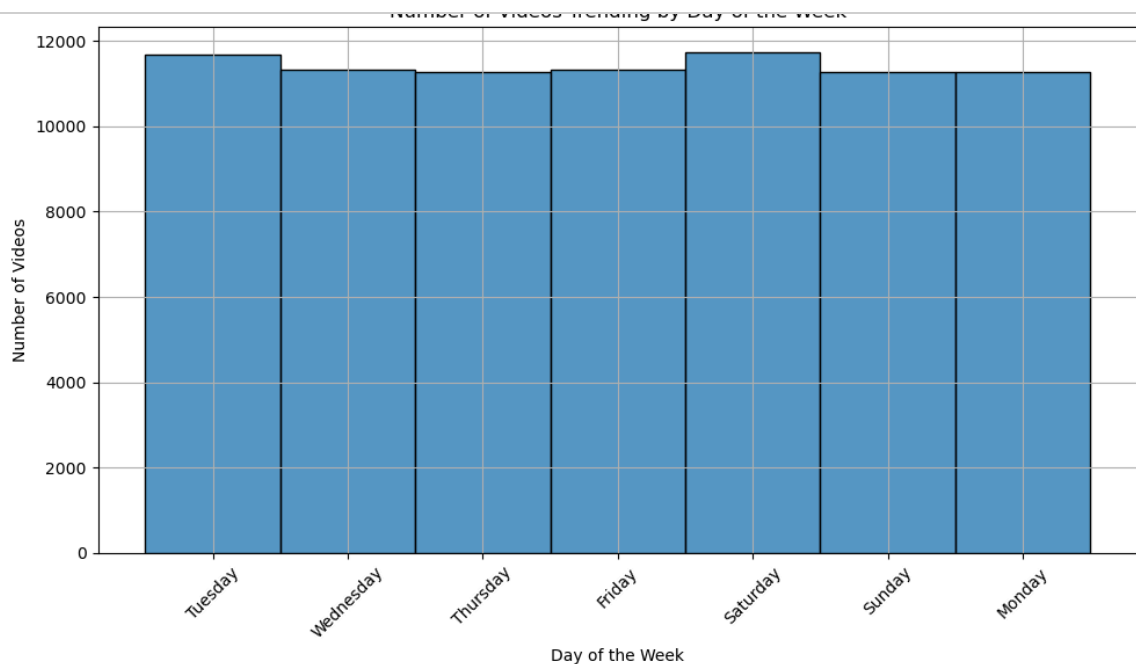
plt.tight_layout()
plt.show()
```



EDA for Date-Time Variables

```
In [8]: # convert the trending_date to datetime type
merged_df['trending_date'] = pd.to_datetime(merged_df['trending_date'])
# Extract day of the week from 'trending_date'
merged_df['trending_day_of_week'] = merged_df['trending_date'].dt.day_of_week

# Plot histogram of trending day of the week
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['trending_day_of_week'], discrete=True, palette='magma')
plt.title('Number of Videos Trending by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Number of Videos')
plt.xticks(rotation=45) # Rotate x-axis labels for better readability
plt.grid(True)
plt.tight_layout()
plt.show()
```




```
In [9]: #convert the publish_date to datetime type
merged_df['publish_time'] = pd.to_datetime(merged_df['publish_time'],
# Extract day of the week from 'publish_time'
merged_df['day_of_week'] = merged_df['publish_time'].dt.day_name()

# Plot histogram of day of the week
plt.figure(figsize=(10, 6))
sns.histplot(merged_df['day_of_week'], discrete=True, palette='viridis')
plt.title('Number of Videos Published by Day of the Week')
plt.xlabel('Day of the Week')
plt.ylabel('Number of Videos')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.show()
```

```
/var/folders/yn/hnpfh1r15tq8t0xq_j4_rzmmh0000gn/T/ipykernel_37366/13
46077495.py:8: UserWarning: Ignoring `palette` because no `hue` var
iable has been assigned.
```

```
    sns.histplot(merged_df['day_of_week'], discrete=True, palette='vi
ridis')
```



Statistical Description

```
In [10]: numerical_description = merged_df.describe()
print(numerical_description)
```

	trending_date	category_id	\
count	79865	79865.000000	
mean	2018-02-25 07:57:45.132410880	18.440205	
min	2017-11-14 00:00:00	1.000000	
25%	2018-01-02 00:00:00	10.000000	
50%	2018-02-23 00:00:00	22.000000	
75%	2018-04-21 00:00:00	24.000000	
max	2018-06-14 00:00:00	43.000000	
std	NaN	7.818304	

	publish_time	views	likes	\
count	79865	7.986500e+04	7.986500e+04	
mean	2018-01-30 08:51:14.599436544	4.091166e+06	1.036262e+05	
min	2006-07-23 08:24:11	5.490000e+02	0.000000e+00	
25%	2017-12-22 15:58:16	2.464170e+05	5.642000e+03	
50%	2018-02-14 05:01:24	7.961060e+05	2.092200e+04	
75%	2018-04-09 08:59:51	2.535704e+06	7.824800e+04	
max	2018-06-14 01:31:53	4.245389e+08	5.613827e+06	
std	NaN	1.439125e+07	2.957265e+05	

```
In [11]: # Statistical description of categorical columns
categorical_description = merged_df[['category_id']].describe()
print(categorical_description)
```

	category_id
count	79865.000000
mean	18.440205
std	7.818304
min	1.000000
25%	10.000000
50%	22.000000
75%	24.000000
max	43.000000

Visualization for Categorical ID


```

In [12]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Group by 'category_id' and count occurrences
category_counts = merged_df.groupby('category_id').size().reset_index()

# Sort by 'N' in descending order
category_counts = category_counts.sort_values(by='N', ascending=False)
category_counts['category_id'] = pd.Categorical(category_counts['category_id'], categories=category_names)

# Create a dictionary to map 'category_id' to descriptive names
category_names = {
    1: "1: Film & Animation",
    2: "2: Autos & Vehicles",
    10: "10: Music",
    15: "15: Pets & Animals",
    17: "17: Sports",
    18: "18: Short Movies",
    19: "19: Travel & Events",
    20: "20: Gaming",
    21: "21: Videoblogging",
    22: "22: People & Blogs",
    23: "23: Comedy",
    24: "24: Entertainment",
    25: "25: News & Politics",
    26: "26: Howto & Style",
    27: "27: Education",
    28: "28: Science & Technology",
    29: "29: Nonprofits & Activism",
    30: "30: Movies",
    31: "31: Anime/Animation",
    32: "32: Action/Adventure",
    33: "33: Classics",
    34: "34: Comedy",
    35: "35: Documentary",
    36: "36: Drama",
    37: "37: Family",
    38: "38: Foreign",
    39: "39: Horror",
    40: "40: Sci-Fi/Fantasy",
    41: "41: Thriller",
    42: "42: Shorts",
    43: "43: Shows",
    44: "44: Trailers"
}

# Map 'category_id' to names in the 'category_counts' DataFrame
category_counts['category_name'] = category_counts['category_id'].map(category_names)

# Plot using seaborn
plt.figure(figsize=(10, 6))
barplot = sns.barplot(data=category_counts, x='category_id', y='N', palette='magma')

# Customize the plot to match your ggplot2 example
plt.title("Top Category ID", fontsize=16)
plt.xlabel(None)

```

```


plt.ylabel(None)
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.figtext(0.9, 0.02, "Donyoe", horizontalalignment='right', fontsize=

# Add a custom legend for category names on the side
handles = barplot.patches
legend_labels = [category_names[int(c)] for c in category_counts['cate

# Position the legend on the right of the plot using 'bbox_to_anchor'
plt.legend(handles=handles[:len(legend_labels)], labels=legend_labels,
           bbox_to_anchor=(1.05, 0.5), loc='center left', borderaxespac

plt.show()

```



Data Transformation-Create Engagement Metrics

```


In [13]: # Create a new column
merged_df['Engagement Metrics'] = merged_df['likes'] + merged_df['dislikes'] + merged_df['comment_count']
# Display the DataFrame to check the new column
print(merged_df[['likes', 'dislikes', 'comment_count', 'Engagement Metrics']])

```

	likes	dislikes	comment_count	Engagement Metrics
0	55681	10247	9479	75407
1	25561	2294	2757	30612
2	787420	43420	125882	956722
3	193	12	37	242
4	30	2	30	62

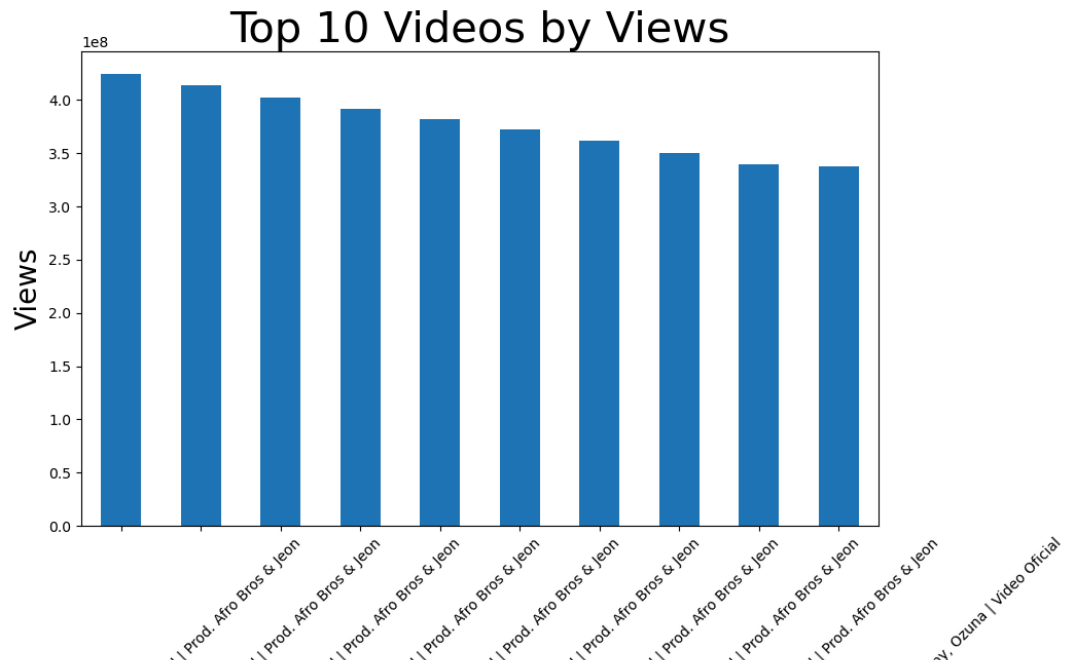
Visualization-Engagement Metrics

```
In [14]: # Create a scatter plot with a regression line
plt.figure(figsize=(8, 6))
sns.regplot(x='Engagement Metrics', y='views', data=merged_df, scatter_
plt.title('Correlation between Engagement Metrics and Views')
plt.xlabel('Engagement Metrics')
plt.ylabel('Views')
plt.show()
```



```
In [15]: top_videos = merged_df.nlargest(10, 'views')[['title', 'views']]
```

```
top_videos.set_index('title')['views'].plot(kind='bar', figsize=(10, 6))
plt.xlabel('Video Title', fontsize=20)
plt.ylabel('Views', fontsize=20)
plt.title('Top 10 Videos by Views', fontsize=30)
plt.xticks(rotation=45)
plt.show()
```



```
In [16]: # Engagement metrics for top 50 videos
top_50_videos = merged_df.nlargest(50, 'views')
print(top_50_videos[['title', 'Engagement Metrics']])
```


```

title Engagement
28412 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
3067426
28212 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
3011515
28008 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2956724
27811 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2902891
27615 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2845332
27424 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2786627
27241 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2723032
27052 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2650114
26861 Nicky Jam x J. Balvin - X (EQUIS) | Video Ofic...
2582210
```

```
In [17]: import seaborn as snb
content = top_50_videos.groupby('channel_title')['views'].max()

# Sort values to get the top 50 channels with the most views
content = content.sort_values(ascending=False).head(50)
content = content.reset_index() # Convert index to column

# Plotting the results
plt.figure(figsize=(14, 8))
snb.barplot(x='channel_title', y='views', data=content)
plt.title('Top 50 Channels with Most views from Top 50 Videos', fontsize=14)
plt.ylabel('views', fontsize=18)
plt.xlabel('Channel', fontsize=18)
plt.xticks(rotation=90)
plt.show()
```



```
In [18]: channel_counts = merged_df.groupby('channel_title')['views'].sum().reset_index()

# Sort values and select top 10 channels
top_10_channels = channel_counts.sort_values(by='views', ascending=False)

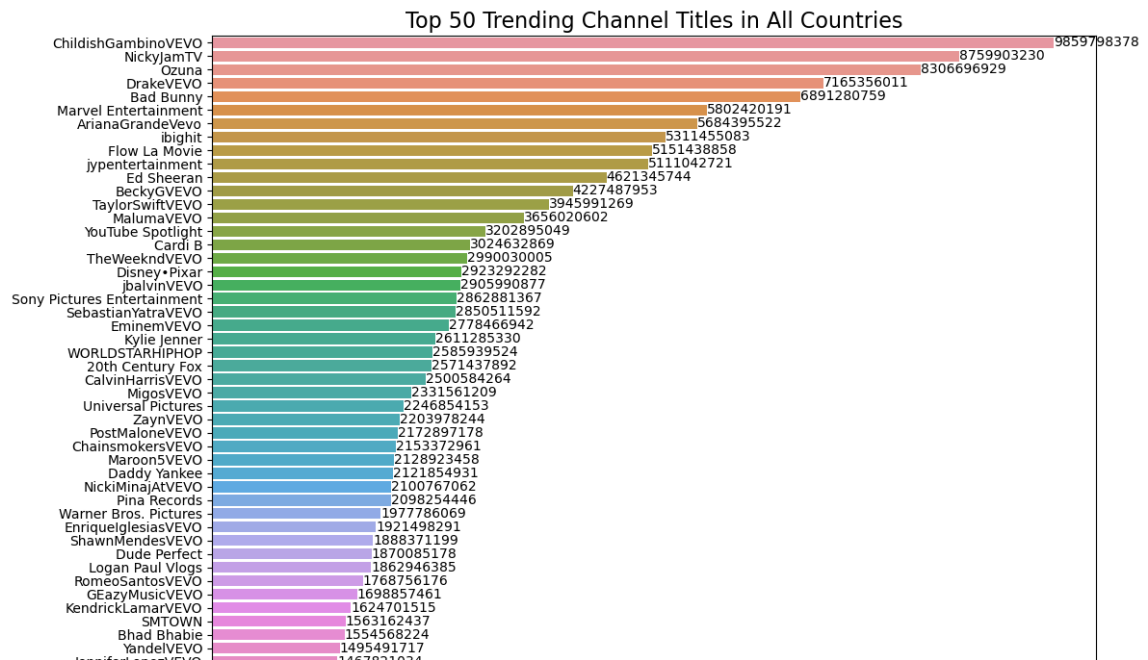
# Plot using seaborn
plt.figure(figsize=(12, 8))
ax = sns.barplot(x='views', y='channel_title', data=top_10_channels, order=top_10_channels['views'].rank())

# Add labels
for index, value in enumerate(top_10_channels['views']):
    ax.text(value, index, str(value), va='center', ha='left', color='black')

# Customize the plot
plt.title('Top 50 Trending Channel Titles in All Countries', fontsize=14)
plt.xlabel('Views', fontsize=12)
plt.ylabel(None)
plt.xticks(rotation=0) # x-axis ticks don't need rotation in horizontal bar chart
plt.tight_layout()

# Add caption
plt.figtext(0.95, 0.02, "Donyoe", horizontalalignment='right', fontsize=10)

# Show the plot
plt.show()
```



Normalize and Standardize Data


Correlation Metrics for Variables

```
In [19]: # add category_id to numerical columns
numerical_columns = ['views', 'likes', 'dislikes', 'comment_count', ]

# Compute the correlation matrix
correlation_matrix = merged_df[numerical_columns].corr()
# Display the correlation matrix
print(correlation_matrix)

# Plot the correlation matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt='.2f',
plt.title('Correlation Heatmap for Numerical Variables')
plt.show()
```

	views	likes	dislikes	comment_count
views	1.000000	0.791670	0.405290	0.485986
likes	0.791670	1.000000	0.448010	0.763192
dislikes	0.405290	0.448010	1.000000	0.745064
comment_count	0.485986	0.763192	0.745064	1.000000



Assign Score for Numerical Values

```
In [20]: import pandas as pd

# Assuming the correlation values are manually entered from the heatmap
correlation_values = {
    'likes': 0.784,          # Correlation of likes with views
    'dislikes': 0.416,      # Correlation of dislikes with views
    'comment_count': 0.502 # Correlation of comment_count with views
}

# Convert the correlation values to absolute values
abs_correlations = {key: abs(value) for key, value in correlation_values.items()}

# Calculate the total sum of absolute correlations
total_correlation = sum(abs_correlations.values())

# Calculate weights by normalizing the absolute correlation values
weights = {key: value / total_correlation for key, value in abs_correlations.items()}

# Convert the weights to a DataFrame for better visualization
weights_df = pd.DataFrame(list(weights.items()), columns=['Variable', 'Weight'])

# Display the weights
print("Calculated Weights of Independent Variables Relative to 'Views':")
print(weights_df)
```

```
Calculated Weights of Independent Variables Relative to 'Views':
   Variable  Weight
0      likes  0.460635
1    dislikes  0.244418
2  comment_count  0.294947
```

```
In [21]: import pandas as pd

weights = {
    'likes': 0.460435,
    'dislikes': 0.244418,
    'comment_count': 0.294947
}

merged_df['score'] = (
    weights['likes'] * merged_df['likes'] -
    weights['dislikes'] * merged_df['dislikes'] +
    weights['comment_count'] * merged_df['comment_count']
)

merged_df['rank'] = merged_df['score'].rank(ascending=False, method='min')

df_sorted = merged_df.sort_values(by='rank')

print(df_sorted)

#output_filename = 'ranked_videos_combined.csv'
#df_sorted.to_csv(output_filename, index=False)

#print("Listing of Every Video with Individual Scores and Ranks Across")
#print(df_sorted[['video_id', 'views', 'likes', 'dislikes', 'comment_count']])
#print(f"\nThe ranking of all videos from all locations has been saved")
```

	video_id	trending_date	title	views	likes	dislikes	comment_count	score	rank
36638	7C2z4GqqS5E	2018-06-01	BTS (방탄소년단) 'FAKE LOVE' Official MV	100000000	10000000	1000000	100000000	100000000	1
77189	7C2z4GqqS5E	2018-06-01	BTS (방탄소년단) 'FAKE LOVE' Official MV	100000000	10000000	1000000	100000000	100000000	2
76988	7C2z4GqqS5E	2018-05-31	BTS (방탄소년단) 'FAKE LOVE' Official MV	100000000	10000000	1000000	100000000	100000000	3
36468	7C2z4GqqS5E	2018-05-31	BTS (방탄소년단) 'FAKE LOVE' Official MV	100000000	10000000	1000000	100000000	100000000	4
36288	7C2z4GqqS5E	2018-05-30	BTS (방탄소년단) 'FAKE LOVE' Official MV	100000000	10000000	1000000	100000000	100000000	5
...
9146	LFhT6H6pRWg	2017-12-29	PSA from Chairman of the FCC Ajit Pai	10000000	1000000	100000	10000000	10000000	1000
9354	LFhT6H6pRWg	2017-12-30	PSA from Chairman of the FCC Ajit Pai	10000000	1000000	100000	10000000	10000000	1001
9575	LFhT6H6pRWg	2017-12-31	PSA from Chairman of the FCC Ajit Pai	10000000	1000000	100000	10000000	10000000	1002

EDA for Score for Top 50 Channels

```
In [22]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

# Assuming your DataFrame is named 'train'
weights = {
    'likes': 0.460435,
    'dislikes': 0.244418,
    'comment_count': 0.294947
}


# Calculate score and rank
merged_df['score'] = (
    weights['likes'] * merged_df['likes'] -
    weights['dislikes'] * merged_df['dislikes'] +
    weights['comment_count'] * merged_df['comment_count']
)

merged_df['rank'] = merged_df['score'].rank(ascending=False, method='min')

# Group by channel_title and sum the scores
channel_scores = merged_df.groupby('channel_title')['score'].sum().reset_index()

# Sort by total score and get top 50 channels
top_channels = channel_scores.sort_values(by='score', ascending=False).head(50)

# Create a bar plot for the top 50 channels
plt.figure(figsize=(12, 8))
sns.barplot(x='score', y='channel_title', data=top_channels, palette='magma')
plt.title('Top 50 Channels by Score')
plt.xlabel('Total Score')
plt.ylabel('Channel Title')
plt.show()
```



Video Titles


```
# Assuming your DataFrame is named 'mergeda_df'
# Concatenate all titles into a single string
all_titles = " ".join(merged_df['title'].astype(str))

# Set up the color palette (equivalent to R's "Dark2")
cmap = Dark2_6.mpl_colormap

# Create a WordCloud object
wordcloud = WordCloud(
    background_color="white",
    max_words=200,
    colormap=cmap,
    width=800,
    height=400,
    random_state=42
)

# Generate the word cloud from the titles
wordcloud.generate(all_titles)

# Plot the word cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off") # Turn off the axis
plt.title('Word Cloud of Video Titles', fontsize=16)
plt.show()
```



Channel Titles


```
In [24]: all_channel_titles = " ".join(merged_df['channel_title'].astype(str))

# Set up the color palette (equivalent to R's "Dark2")
cmap = Dark2_6.mpl_colormap

# Create a WordCloud object
wordcloud = WordCloud(
    background_color="white",
    max_words=200,
    colormap=cmap,
    width=800,
    height=400,
    random_state=42
)

# Generate the word cloud from the titles
wordcloud.generate(all_channel_titles)

# Plot the word cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off") # Turn off the axis
plt.title('Word Cloud of Channel Titles', fontsize=16)
plt.show()
```




```
# Creating Word Cloud-tags
all_tags = " ".join(merged_df['tags'].astype(str))

# Set up the color palette (equivalent to R's "Dark2")
cmap = Dark2_6.mpl_colormap

# Create a WordCloud object
wordcloud = WordCloud(
    background_color="white",
    max_words=200,
    colormap=cmap,
    width=800,
    height=400,
    random_state=42
)

# Generate the word cloud from the titles
wordcloud.generate(all_tags)

# Plot the word cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off") # Turn off the axis
plt.title('Word Cloud of Tags', fontsize=16)
plt.show()
```



Video Descriptions


```
In [26]: all_description = " ".join(merged_df['description'].astype(str))

# Set up the color palette (equivalent to R's "Dark2")
cmap = Dark2_6.mpl_colormap

# Create a WordCloud object
wordcloud = WordCloud(
    background_color="white",
    max_words=200,
    colormap=cmap,
    width=800,
    height=400,
    random_state=42
)

# Generate the word cloud from the titles
wordcloud.generate(all_description)

# Plot the word cloud
plt.figure(figsize=(10, 6))
plt.imshow(wordcloud, interpolation="bilinear")
plt.axis("off") # Turn off the axis
plt.title('Word Cloud of Video Descriptions', fontsize=16)
plt.show()
```



Drop Unnecessary Columns

```
In [27]: #drop columns needed
merged_df.drop(columns=['thumbnail_link', 'video_id', 'comments_disabled'])
print(merged_df.head())
```

		tags	views	li
0	christmas "john lewis christmas" "john lewis" ...		7224515	55
1	SNL "Saturday Night Live" "SNL Season 43" "Epi...		1053632	25
2	Eminem "Walk" "On" "Water" "Aftermath/Shady/In...		17158579	787
3	Salford City FC "Salford City" "Salford" "Clas...		27833	
4		[none]	9815	

	dislikes	comment_count	d
0	10247	9479	Click here to continue the story and ma
1	2294	2757	Musical guest Taylor Swift performs ...Re
2	12422	125002	Eminem's new track Walk on Water ft. D-

Text Preprocessing


```
In [28]: import nltk
nltk.download('stopwords')
from nltk.corpus import stopwords
import re

# Get the list of default English stopwords
stop_words = set(stopwords.words('english'))

# Function to remove stopwords and clean text
def clean_text(text):
    # Lowercase the text
    text = text.lower()

    # Remove non-alphabetical characters (retain only letters and space)
    text = re.sub(r'^a-z\s', '', text)

    # Split text into words
    words = text.split()

    # Remove stopwords
    remove_stopwords = [word for word in words if word not in stop_words]

    # Join the cleaned words back into a string
    new_text = ' '.join(remove_stopwords)

    return new_text
data = {'title', 'description', 'text'}
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] /Users/yuhanzhao/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```

      trending_date                                ti
title \
0      2017-11-14      John Lewis Christmas Ad 2017 - #MozTheMons
ter
1      2017-11-14      Taylor Swift: ...Ready for It? (Live) -
SNL
2      2017-11-14      Eminem - Walk On Water (Audio) ft. Beyo
ncé
3      2017-11-14      Goals from Salford City vs Class of 92 and Fr
i...
4      2017-11-14      Dashcam captures truck's near miss with child
...
...      ...
...
79860  2018-06-14      The Cat Who Caught the La
ser
```

```
In [29]: # Check the data types of each column
print(merged_df.dtypes)
```

```
trending_date      datetime64[ns]
title              object
channel_title      object
category_id        int64
publish_time       datetime64[ns]
tags              object
views             int64
likes             int64
dislikes          int64
comment_count      int64
description        object
trending_day_of_week object
day_of_week        object
Engagement Metrics int64
score             float64
rank              float64
new_text          object
dtype: object
```

Split the Dataset into Train and Test by 80/20

```
In [30]: from sklearn.model_selection import train_test_split

X = merged_df.drop(columns=['views']) # Drop 'views' from features to
y = merged_df['views']
# Assuming you have a dataset with features X and target y
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size=0.2,

train = pd.DataFrame(X_train)
train['views'] = y_train.values

test = pd.DataFrame(X_test)
test['views'] = y_test.values
```

Feature Engineering


Days Since Published

```
In [31]: #convert the type of publish time
train['publish_time'] = pd.to_datetime(train['publish_time'])
train['trending_date'] = pd.to_datetime(train['trending_date'], format='%Y-%m-%d')

# Creating a new feature 'days_since_published'
train['days_since_published'] = (train['trending_date'] - train['publish_time']).dt.days

# Creating bins for days since published
bins = [0, 7, 14, 30, 60, 90, 120, 180, 365] # Example bins
labels = ['0-7', '8-14', '15-30', '31-60', '61-90', '91-120', '121-180', '181-365']
train['days_bins'] = pd.cut(train['days_since_published'], bins=bins, labels=labels)

plt.figure(figsize=(10, 6))
sns.boxplot(data=train, x='days_bins', y='views')
plt.title('Box Plot of Views by Days Since Published Bins')
plt.xlabel('Days Since Published Bins')
plt.ylabel('View Count')
plt.grid(True)
plt.show()
```



Sentimental Analysis

Sentiment Polarity Distribution

```
In [32]: from textblob import TextBlob
import matplotlib.pyplot as plt

# Calculate sentiment polarity for description and title
def get_sentiment(text):
    return TextBlob(text).sentiment.polarity

# Apply sentiment analysis
train['description_sentiment'] = train['description'].fillna('').apply(get_sentiment)
train['title_sentiment'] = train['title'].fillna('').apply(get_sentiment)

# Calculate average sentiment scores
avg_description_sentiment = train['description_sentiment'].mean()
avg_title_sentiment = train['title_sentiment'].mean()

print("Average Description Sentiment Score:", avg_description_sentiment)
print("Average Title Sentiment Score:", avg_title_sentiment)

# Plotting the sentiment distributions
plt.figure(figsize=(14, 6))

# Description Sentiment Histogram
plt.subplot(1, 2, 1)
plt.hist(train['description_sentiment'], bins=30, alpha=0.7, edgecolor='black')
plt.title('Description Sentiment Polarity Distribution')
plt.xlabel('Sentiment Polarity')
plt.ylabel('Frequency')
plt.xticks([-1, 0, 1])

# Title Sentiment Histogram
plt.subplot(1, 2, 2)
plt.hist(train['title_sentiment'], bins=30, alpha=0.7, edgecolor='black')
plt.title('Title Sentiment Polarity Distribution')
plt.xlabel('Sentiment Polarity')
plt.ylabel('Frequency')
plt.xticks([-1, 0, 1])

plt.tight_layout()
plt.show()
```

Average Description Sentiment Score: 0.1716764242965884
Average Title Sentiment Score: 0.0477964529239135



Visualize the Sentiment Distrubution Category

```
In [33]: import pandas as pd
import matplotlib.pyplot as plt

# Define sentiment categories
def categorize_sentiment(polarity):
    if polarity > 0:
        return 'Positive'
    elif polarity < 0:
        return 'Negative'
    else:
        return 'Neutral'

# Apply categorization to sentiment columns
train['description_sentiment_category'] = train['description_sentiment']
train['title_sentiment_category'] = train['title_sentiment'].apply(cate

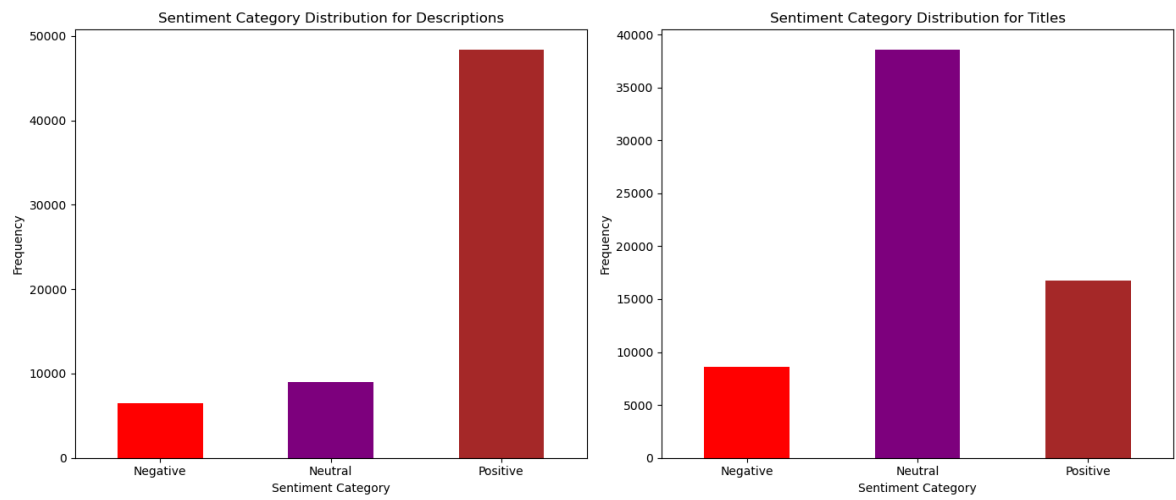
# Plot sentiment category distribution for descriptions and titles
plt.figure(figsize=(14, 6))

# Custom order for categories
category_order = ['Negative', 'Neutral', 'Positive']

# Plot description sentiment distribution
plt.subplot(1, 2, 1)
description_sentiment_counts = train['description_sentiment_category'].value_counts()
description_sentiment_counts.plot(kind='bar', color=['red', 'purple', 'brown'])
plt.title('Sentiment Category Distribution for Descriptions')
plt.xlabel('Sentiment Category')
plt.ylabel('Frequency')
plt.xticks(rotation=0)

# Plot title sentiment distribution
plt.subplot(1, 2, 2)
title_sentiment_counts = train['title_sentiment_category'].value_counts()
title_sentiment_counts.plot(kind='bar', color=['red', 'purple', 'brown'])
plt.title('Sentiment Category Distribution for Titles')
plt.xlabel('Sentiment Category')
plt.ylabel('Frequency')
plt.xticks(rotation=0)

plt.tight_layout()
plt.show()
```



Create TF-IDF Feature

description Column

```

In [34]: from sklearn.feature_extraction.text import TfidfVectorizer
import numpy as np

# Ensure the 'description' column exists in the DataFrame
if 'description' in train.columns:
    # Assuming 'description' column contains the text data
    text_data = train['description'].fillna('') # Handle missing values

    # Check if text_data is iterable, not a single string
    if isinstance(text_data, pd.Series):
        # Initialize the TF-IDF Vectorizer
        tfidf_vectorizer = TfidfVectorizer(max_features=100, stop_words=

        # Fit and transform the text data to generate the TF-IDF matrix
        tfidf_matrix = tfidf_vectorizer.fit_transform(text_data)

        # Convert the sparse matrix into a DataFrame for easier manipulation
        tfidf_df = pd.DataFrame(tfidf_matrix.toarray(), columns=tfidf_v

        # Function to get top N features per row based on TF-IDF score
        def get_top_tfidf_features(row, features, top_n=5):
            top_indices = np.argsort(row)[-1][:top_n] # Get the indices
            top_features = [(features[i], row[i]) for i in top_indices]
            return top_features

        # Apply the function to each row in the TF-IDF matrix
        top_tfidf_features = [get_top_tfidf_features(row, tfidf_vectorizer.get_feature_names(), 5)
                               for row in tfidf_matrix.toarray()]

        # Add the top TF-IDF features as a new column in the original DataFrame
        train['top_tfidf_features'] = top_tfidf_features

        # Display the entire first 5 rows of the DataFrame including the new column
        print(train.head(5))
    else:
        print("The 'description' column should be a pandas Series.")
else:
    print("The DataFrame does not contain a 'description' column.")

```


	trending_date	title
23604	2018-03-14	Marshmello & Anne-Marie: Friends
25630	2018-03-24	Kirby Star Allies' Surprising HD Rumble Secret...
68698	2018-04-20	Stephen A.: Kevin Hart 'got his feelings hurt'...
39559	2017-11-17	How to be an Aquarius
62877	2018-03-16	Charlie Puth - Done For Me (feat. Kehlani) [Official Music Video]

	channel_title	category_id
23604	The Tonight Show Starring Jimmy Fallon	23
25630	GameXplain	20
68698	ESPN	17
39559	Sailor J	24
62877	Charlie Puth	10

	publish_time	tags
23604	2018-03-07 14:00:03	The Tonight Show Jimmy Fallon Marshmello ...
25630	2018-03-16 04:00:01	Kirby Kirby Star Allies Dedede Meta Knight ...
68698	2018-04-17 14:55:31	espn dwyanedwade dwyanedwade dwyanedwade ...
39559	2017-11-15 13:29:28	Zodiac makeup comedy aquarius ...
62877	2018-03-15 16:02:17	Charlie Puth charlieputh CharliePuth - Done For Me ...

	likes	dislikes	comment_count
23604	45011	1156	2365
25630	2716	52	450
68698	6829	537	1445
39559	5172	453	976
62877	84227	739	8663

	description	rank
23604	Music guest Marshmello & Anne-Marie performs Friends	27872
25630	Kirby Star Allies does something pretty fun with	67076
68698	First Take's Stephen A. Smith says Kevin Hart is	57168
39559	Ya'll asked lol. What sign should I do next? Darius	60832
62877	Download & Stream Done For Me (feat. Kehlani): Official Music Video	18947

	new_text	views
23604	marshmello annemarie friends	1443792
25630	kirby star allies surprising hd rumble secret	106398
68698	stephen kevin hart got feelings hurt dwyanedwade	976783

```

39559                                aquarius      88644
62877    charlie puth done feat kehlani official audio  722009

```

```

      days_since_published  days_bins  description_sentiment  title_s
entiment \
23604                6      0-7                0.232292
0.000000
25630                7      0-7                0.137500
0.141667
68698                2      0-7                0.168333
0.250000
39559                1      0-7                0.400000
0.000000
62877                0      NaN                0.400000
0.000000

```

```

      description_sentiment_category  title_sentiment_category \
23604                Positive                Neutral
25630                Positive                Positive
68698                Positive                Positive
39559                Positive                Neutral
62877                Positive                Neutral

```

```

      top_tfidf_features
23604  [(jimmy, 0.6973441834478303), (nbc, 0.47929685...
25630  [(patreon, 0.5185912254067347), (com, 0.407894...
68698  [(http, 0.6127313163416526), (youtube, 0.33885...
39559  [(ll, 0.5094338648331312), (don, 0.47521836732...
62877  [(nhttp, 0.4915874738421169), (com, 0.44499439...

```

[5 rows x 24 columns]

tags Column

```

In [35]: from sklearn.feature_extraction.text import TfidfVectorizer
import numpy as np

# Ensure the 'description' column exists in the DataFrame
if 'tags' in train.columns:
    # Assuming 'description' column contains the text data
    text_data = train['tags'].fillna('') # Handle missing values

    # Check if text_data is iterable, not a single string
    if isinstance(text_data, pd.Series):
        # Initialize the TF-IDF Vectorizer
        tfidf_vectorizer = TfidfVectorizer(max_features=100, stop_words=

        # Fit and transform the text data to generate the TF-IDF matrix
        tfidf_matrix = tfidf_vectorizer.fit_transform(text_data)

        # Convert the sparse matrix into a DataFrame for easier manipulation
        tfidf_df = pd.DataFrame(tfidf_matrix.toarray(), columns=tfidf_v

        # Function to get top N features per row based on TF-IDF score
        def get_top_tfidf_features(row, features, top_n=5):
            top_indices = np.argsort(row)[-1][:top_n] # Get the indices
            top_features = [(features[i], row[i]) for i in top_indices]
            return top_features

        # Apply the function to each row in the TF-IDF matrix
        top_tfidf_features = [get_top_tfidf_features(row, tfidf_vectorizer.get_feature_names(), 5)
                               for row in tfidf_matrix.toarray()]

        # Add the top TF-IDF features as a new column in the original DataFrame
        train['top_tfidf_features'] = top_tfidf_features

        # Display the entire first 5 rows of the DataFrame including the new column
        print(train.head(5))
    else:
        print("The 'description' column should be a pandas Series.")
else:
    print("The DataFrame does not contain a 'description' column.")

```

	trending_date	ti
tle \		
23604	2018-03-14	Marshmello & Anne-Marie: Frie
nds		
25630	2018-03-24	Kirby Star Allies' Surprising HD Rumble Secre
t...		
68698	2018-04-20	Stephen A.: Kevin Hart 'got his feelings hur
t'...		
39559	2017-11-17	How to be an Aquar
ius		
62877	2018-03-16	Charlie Puth - Done For Me (feat. Kehlani) [0
f...		

	channel_title	category_id	\
23604	The Tonight Show Starring Jimmy Fallon	23	
25630	GameXplain	20	
68698	ESPN	17	
39559	Sailor J	24	
62877	Charlie Puth	10	

Dimension Reduction-PCA

```
In [36]: from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA

non_numeric_cols = ['publish_time', 'title', 'channel_title', 'tags',
X_train_model = train.drop(columns=non_numeric_cols + ['views']).select_
X_test_model = test.drop(columns=non_numeric_cols + ['views']).select_

X_test_model = X_test_model.reindex(columns=X_train_model.columns, fill

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_model)
X_test_scaled = scaler.transform(X_test_model)

print("Missing values in X_train_model:\n", X_train_model.isna().sum())
print("Missing values in X_test_model:\n", X_test_model.isna().sum())
```

Missing values in X_train_model:

category_id	0
likes	0
dislikes	0
comment_count	0
Engagement Metrics	0
score	0
rank	0
days_since_published	0
description_sentiment	0
title_sentiment	0
dtype: int64	


Missing values in X_test_model:

category_id	0
likes	0
dislikes	0
comment_count	0
Engagement Metrics	0
score	0
rank	0

```
In [37]: # Apply PCA (Reduce to n components to capture 95% of variance)
pca = PCA(n_components=0.95)
X_train_pca = pca.fit_transform(X_train_scaled)
X_test_pca = pca.transform(X_test_scaled)

# Visualize the PCA results (Plot only the first two components)
plt.figure(figsize=(10, 6))
plt.scatter(X_train_pca[:, 0], X_train_pca[:, 1], c=y_train, cmap='viridis')
plt.colorbar(label='Views')
plt.title('PCA of YouTube Data (Train Set)', weight='bold')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.grid(True)
plt.show()

# Explained variance for all components selected by PCA
explained_variance = pca.explained_variance_ratio_
print("Explained Variance per component:")
for i, variance in enumerate(explained_variance, start=1):
    print(f"PC{i}: {variance:.2%}")
```



Model Building

```
In [38]: import xgboost as xgb
from sklearn.metrics import mean_squared_error, r2_score
import numpy as np
import pandas as pd
```

```
In [39]: # Define a function for calculating model metrics
def calculate_metrics(model, X_train, y_train, X_test, y_test):
    train_preds = model.predict(X_train)
    test_preds = model.predict(X_test)

    # Calculate RMSE and R^2 for training and test sets
    train_rmse = np.sqrt(mean_squared_error(y_train, train_preds))
    test_rmse = np.sqrt(mean_squared_error(y_test, test_preds))
    train_r2 = r2_score(y_train, train_preds)
    test_r2 = r2_score(y_test, test_preds)

    return {
        "Train RMSE": train_rmse, "Test RMSE": test_rmse,
        "Train R^2": train_r2, "Test R^2": test_r2
    }

In [40]: # Define a function to train the model with specific hyperparameters
def train_xgboost(X_train, y_train, X_test, y_test, params):
    model = xgb.XGBRegressor(**params, random_state=42)
    model.fit(X_train, y_train)

    # Calculate and return metrics
    metrics = calculate_metrics(model, X_train, y_train, X_test, y_test)
    return model, metrics

In [41]: # Define hyperparameter variations
variations = [
    {"learning_rate": 0.05, "n_estimators": 200, "max_depth": 6}
]

In [42]: # Initialize a DataFrame to store results for each variation
results = pd.DataFrame(columns=["Variation", "Train RMSE", "Test RMSE"])
```

```

In [43]: # Create a list of columns to drop if they exist
text_columns = ['title', 'channel_title', 'tags', 'description', 'location']
X_train.drop([col for col in text_columns if col in X_train.columns], axis=1, inplace=True)
X_test.drop([col for col in text_columns if col in X_test.columns], axis=1, inplace=True)

# Encode categorical features using one-hot encoding for consistency
categorical_columns = ['trending_day_of_week', 'day_of_week', 'days_buried']
X_train = pd.get_dummies(X_train, columns=[col for col in categorical_columns])
X_test = pd.get_dummies(X_test, columns=[col for col in categorical_columns])

# Convert datetime columns to relevant features if they exist
if 'trending_date' in X_train.columns:
    X_train['trending_year'] = X_train['trending_date'].dt.year
    X_train['trending_month'] = X_train['trending_date'].dt.month
    X_train['trending_day'] = X_train['trending_date'].dt.day
    X_train.drop(['trending_date'], axis=1, inplace=True)

if 'trending_date' in X_test.columns:
    X_test['trending_year'] = X_test['trending_date'].dt.year
    X_test['trending_month'] = X_test['trending_date'].dt.month
    X_test['trending_day'] = X_test['trending_date'].dt.day
    X_test.drop(['trending_date'], axis=1, inplace=True)

# Drop 'publish_time' if it exists
if 'publish_time' in X_train.columns:
    X_train.drop(['publish_time'], axis=1, inplace=True)

if 'publish_time' in X_test.columns:
    X_test.drop(['publish_time'], axis=1, inplace=True)

# Ensure X_test has the same columns as X_train
X_test = X_test.reindex(columns=X_train.columns, fill_value=0)

# Verify that X_train and X_test now have the same columns
print("X_train columns:", X_train.columns)
print("X_test columns:", X_test.columns)

```



```

X_train columns: Index(['category_id', 'likes', 'dislikes', 'comment_count',
                        'Engagement Metrics', 'score', 'rank', 'trending_day_of_week_Monday',
                        'trending_day_of_week_Saturday', 'trending_day_of_week_Sunday',
                        'trending_day_of_week_Thursday', 'trending_day_of_week_Tuesday',
                        'trending_day_of_week_Wednesday', 'day_of_week_Monday',
                        'day_of_week_Saturday', 'day_of_week_Sunday', 'day_of_week_Thursday',
                        'day_of_week_Tuesday', 'day_of_week_Wednesday', 'trending_year',
                        'trending_month', 'trending_day'],
                        dtype='object')
X_test columns: Index(['category_id', 'likes', 'dislikes', 'comment_count',
                       'Engagement Metrics', 'score', 'rank', 'trending_day_of_week_Monday',
                       'trending_day_of_week_Saturday', 'trending_day_of_week_Sunday',
                       'trending_day_of_week_Thursday', 'trending_day_of_week_Tuesday',
                       'trending_day_of_week_Wednesday', 'day_of_week_Monday',
                       'day_of_week_Saturday', 'day_of_week_Sunday', 'day_of_week_Thursday',
                       'day_of_week_Tuesday', 'day_of_week_Wednesday', 'trending_year',
                       'trending_month', 'trending_day'],
                       dtype='object')

```

```

In [44]: def calculate_metrics(model, X_train, y_train, X_test, y_test):
          # Predictions
          train_preds = model.predict(X_train)
          test_preds = model.predict(X_test)

          # Calculate metrics
          train_rmse = np.sqrt(mean_squared_error(y_train, train_preds))
          test_rmse = np.sqrt(mean_squared_error(y_test, test_preds))

          train_r2 = r2_score(y_train, train_preds)
          test_r2 = r2_score(y_test, test_preds)

          return {
              "Train RMSE": train_rmse,
              "Test RMSE": test_rmse, # Changed from val_rmse to test_rmse
              "Train R^2": train_r2,
              "Test R^2": test_r2
          }

```

```
In [45]: # Create an empty DataFrame if it isn't already
results = pd.DataFrame()

# Train models for each variation and record results
for i, params in enumerate(variations):
    model, metrics = train_xgboost(X_train, y_train, X_test, y_test, pa

    # Create a DataFrame with the metrics for this variation
    result_row = pd.DataFrame({
        "Variation": [f"Variation {i + 1}"],
        **metrics
    })

    # Concatenate the new row to the results DataFrame
    results = pd.concat([results, result_row], ignore_index=True)
```

```
In [46]: # Display the comparison table
print("Comparison of XGBoost Model Variations:")
print(results)

# Identify the best model based on Validation RMSE
best_model_index = results["Test RMSE"].idxmin()
best_params = variations[best_model_index]
print(f"\nBest Model Variation: {best_model_index + 1}")
print(f"Hyperparameters: {best_params}")
print(results.iloc[best_model_index])
```

Comparison of XGBoost Model Variations:

	Variation	Train RMSE	Test RMSE	Train R ²	Test R ²
0	Variation 1	2.120971e+06	2.909246e+06	0.978047	0.960791

Best Model Variation: 1

Hyperparameters: {'learning_rate': 0.05, 'n_estimators': 200, 'max_depth': 6}

	Variation 1
Train RMSE	2120971.173631
Test RMSE	2909245.576456
Train R ²	0.978047
Test R ²	0.960791

Name: 0, dtype: object

Explain the model

Feature Importance

```
In [47]: # Get feature importances directly
feature_importances = model.feature_importances_

# Adjust to take only the first 17 feature importances
feature_importances_adjusted = feature_importances[:len(X.columns)]

# Generate feature importance DataFrame
top_features = pd.DataFrame({'Feature': X.columns, 'Importance': feature_importances_adjusted})
top_10_features = top_features.head(10)

print("Top 10 Features:\n", top_10_features)
```

Top 10 Features:

	Feature	Importance
1	title	0.437080
4	publish_time	0.193274
2	channel_title	0.085307
5	tags	0.051353
3	category_id	0.048652
14	rank	0.040345
15	new_text	0.027795
6	likes	0.009272
13	score	0.009106
0	trending_date	0.008889

Extracting and Analyzing 5 Individual Predictions

```
In [48]: import shap
import numpy as np

# Randomly select 5 samples from the test set
random_samples = X_test.sample(5, random_state=42)

explainer = shap.TreeExplainer(model)
shap_values = explainer.shap_values(random_samples)

# Generate and analyze SHAP values for random samples
for i, sample_index in enumerate(random_samples.index):
    print(f"\nExplanation for sample {i+1} (Index: {sample_index}):")

    # Display all columns' information for this prediction
    print("Sample details:")
    display(random_samples.loc[[sample_index]]) # Show the full row details

    # Calculate and display the SHAP force plot for visualization
    shap.initjs()
    shap.force_plot(explainer.expected_value, shap_values[i], random_samples.loc[[sample_index]])

    # Retrieve SHAP values for the sample and sort by absolute importance
    sample_shap_values = shap_values[i]
    feature_importances = pd.DataFrame({
        'feature': random_samples.columns,
        'shap_value': sample_shap_values
    }).set_index('feature').sort_values(by='shap_value', key=abs, ascending=False)

    # Display the top features that influenced the prediction
    print("Top contributing features:")
    print(feature_importances.head())

    # Suggest changes for flipping the prediction
    current_prediction = model.predict(random_samples.iloc[[i]])[0]
    if current_prediction == 1:
        print(f"To flip from 1 to 0, consider reducing the values of the top contributing features:")
    else:
        print(f"To flip from 0 to 1, consider increasing the values of the top contributing features:")

    # Suggest how to change the feature values
    for feature, shap_value in feature_importances.head().itertuples():
        direction = "increase" if shap_value < 0 else "decrease"
        print(f" - {feature}: Consider a {direction} of {abs(shap_value)} to flip the prediction")
```

Explanation for sample 1 (Index: 49053):
Sample details:

	category_id	likes	dislikes	comment_count	Engagement Metrics	score	rank	trend
	49053	23	2837	86	191	3114	1341.569024	67271.0

1 rows x 22 columns



higher ↕ lower

Model Deployment, Data Drift and Concept Drift, and Model Monitoring

Serializing the Model with Pickle

```
In [50]: import pickle

model_filename = 'XGBoost_model.pkl'

# Save the trained model to a file
with open(model_filename, 'wb') as file:
    pickle.dump(model, file)

print(f"Model saved to {model_filename}")
```

Model saved to XGBoost_model.pkl

```
In [51]: # Load the model from the file
with open(model_filename, 'rb') as file:
    loaded_model = pickle.load(file)
```

```
In [52]: predictions = loaded_model.predict(X_test)
```

Save and Load the Data

```
In [53]: data_filename = 'data.pkl'

# Save datasets to a file
with open(data_filename, 'wb') as file:
    pickle.dump({'X_train': X_train, 'X_test': X_test, 'y_train': y_train, 'y_test': y_test}, file)

print(f"Datasets saved to {data_filename}")
```

Datasets saved to data.pkl

```
In [54]: # Load datasets from the file
with open(data_filename, 'rb') as file:
    data = pickle.load(file)

# Extract datasets
X_train = data['X_train']
X_test = data['X_test']
y_train = data['y_train']
y_test = data['y_test']
```

```
In [55]: import platform
import sys
import pkg_resources

# Collect OS and Python version information
print("OS:", platform.system(), platform.version())
print(f"Operating System: {platform.system()} {platform.release()}")
print("Python Version:", sys.version)
```

```
OS: Darwin Darwin Kernel Version 23.3.0: Wed Dec 20 21:28:58 PST 202
3; root:xnu-10002.81.5~7/RELEASE_X86_64
Operating System: Darwin 23.3.0
Python Version: 3.11.5 (main, Sep 11 2023, 08:19:27) [Clang 14.0.6 ]
```

```
In [56]: import importlib
import pkg_resources

# List of packages identified from notebook
packages = [
    "matplotlib.pyplot", "nltk", "numpy", "os", "palettable.colorbrewer",
    "pandas", "re", "seaborn", "shap", "sklearn.decomposition",
    "sklearn.feature_extraction.text", "sklearn.metrics", "sklearn.model_selection",
    "sklearn.preprocessing", "textblob", "wordcloud", "xgboost", "zipfile"
]

# Function to get package version
def get_package_version(package):
    try:
        # Handle "matplotlib.pyplot" style imports
        if '.' in package:
            package = package.split('.')[0]
        # Try pkg_resources first
        version = pkg_resources.get_distribution(package).version
    except pkg_resources.DistributionNotFound:
        # If not found, try importlib
        try:
            module = importlib.import_module(package)
            version = getattr(module, '__version__', 'Version not found')
        except ImportError:
            version = 'Not installed'
    return version

# Collect versions
package_versions = {pkg: get_package_version(pkg) for pkg in packages}

# Display package versions
print("Package Versions:")
for pkg, version in package_versions.items():
    print(f"{pkg}: {version}")
```

```
Package Versions:
matplotlib.pyplot: 3.7.2
nltk: 3.8.1
numpy: 1.24.3
os: Version not found
palettable.colorbrewer.qualitative: 3.3.3
pandas: 2.0.3
re: 2.2.1
seaborn: 0.12.2
shap: 0.46.0
sklearn.decomposition: 1.3.0
sklearn.feature_extraction.text: 1.3.0
sklearn.metrics: 1.3.0
sklearn.model_selection: 1.3.0
sklearn.preprocessing: 1.3.0
textblob: 0.18.0.post0
wordcloud: 1.9.3
xgboost: 2.0.3
zipfile: Version not found
```

